

# Image Classification w/ DL

Lets import all the neccessary libraries and set some variables

```
In [30]: #All Imports
import os
import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
import seaborn as sb
import tensorflow as tf
import keras
from keras.applications import InceptionV3

# Unzip after uploading
# %reset
# !unzip '/content/cards/data.zip'
# !rm -rf '/content/__MACOSX'

#Constants:
batch_size = 40
image_size = (224,224)
num_classes = 4
epochs = 10
```

## Create Sets

We are going to try to compare models in their ability to classify playing cards by suit

So we start by creating the train and test sets, and show a few examples of the types of images in the set.

```
In [26]: #Create train and test sets
train_data, test_data = tf.keras.utils.image_dataset_from_directory('/content')

#Check the inferred classes
class_names = train_data.class_names
print(class_names)

#Found at [https://www.tensorflow.org/tutorials/images/transfer_learning]
plt.figure(figsize=(10, 10))
for images, labels in train_data.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(images[i].numpy().astype("uint8"))
        #print(list((labels[i].numpy().astype('int'))).index(1))
        plt.title(class_names[list((labels[i].numpy().astype('int'))).index(1)])
        plt.axis("off")

train_data = train_data.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
test_data = test_data.cache().prefetch(buffer_size=tf.data.AUTOTUNE)
```

```
Found 8024 files belonging to 4 classes.
Using 6420 files for training.
Using 1604 files for validation.
['Club', 'Diamond', 'Heart', 'Spade']
```

Heart



Diamond




Club




Heart



Club



Club



Diamond



Spade



Spade



## Sequential Model

This is the a normal Sequential Model

```
In [4]: sequential = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(224,224,3)),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])

sequential.compile(
    loss = 'categorical_crossentropy',
    optimizer = 'rmsprop',
    metrics = ['accuracy']
)

sequential.summary()

history_sequential = sequential.fit(
    train_data,
    epochs = epochs,
    steps_per_epoch = 15,
    batch_size = batch_size,
    validation_data = test_data
)
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 150528)	0
dense_3 (Dense)	(None, 512)	77070848
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 512)	262656
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 4)	2052

=====  
 Total params: 77,335,556  
 Trainable params: 77,335,556  
 Non-trainable params: 0

Epoch 1/10

15/15 [=====] - 25s 2s/step - loss: 69479.0234 - accuracy: 0.2550 - val\_loss: 16944.3984 - val\_accuracy: 0.2450

Epoch 2/10

15/15 [=====] - 25s 2s/step - loss: 5222.2573 - accuracy: 0.2350 - val\_loss: 13.5702 - val\_accuracy: 0.2531

Epoch 3/10

15/15 [=====] - 24s 2s/step - loss: 28.7262 - accuracy: 0.2155 - val\_loss: 3.1651 - val\_accuracy: 0.2481

Epoch 4/10

15/15 [=====] - 28s 2s/step - loss: 6.2252 - accuracy: 0.2250 - val\_loss: 2.9856 - val\_accuracy: 0.2419

Epoch 5/10

15/15 [=====] - 27s 2s/step - loss: 2.5171 - accuracy: 0.2600 - val\_loss: 4.1997 - val\_accuracy: 0.2712

Epoch 6/10

15/15 [=====] - 25s 2s/step - loss: 18.6030 - accuracy: 0.2533 - val\_loss: 4.5738 - val\_accuracy: 0.2687

Epoch 7/10

15/15 [=====] - 28s 2s/step - loss: 3.3126 - accuracy: 0.3033 - val\_loss: 2.8934 - val\_accuracy: 0.2693

Epoch 8/10

15/15 [=====] - 30s 2s/step - loss: 8.0784 - accuracy: 0.2700 - val\_loss: 2.4039 - val\_accuracy: 0.2681


Epoch 9/10

15/15 [=====] - 29s 2s/step - loss: 1.6194 - accuracy: 0.3000 - val\_loss: 2.3869 - val\_accuracy: 0.2693

Epoch 10/10

15/15 [=====] - 25s 2s/step - loss: 38.6142 - accuracy: 0.2500 - val\_loss: 1.4601 - val\_accuracy: 0.2712

```
In [5]: plt.plot(history_sequential.history['val_accuracy'])  
plt.plot(history_sequential.history['accuracy'])  
plt.title('Sequential Accuracy')  
plt.legend(['Train', 'Test'])  
plt.show()
```



The fact that accuracy goes up and down means the learning rate is too high.

## CNN Model

```
In [8]: cnn = tf.keras.models.Sequential([
    tf.keras.Input(shape=(224,224,3)),
    tf.keras.layers.Conv2D(32, kernel_size=(3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2)),
    tf.keras.layers.Conv2D(32, kernel_size=(3,3), activation='relu'),
    tf.keras.layers.MaxPooling2D(pool_size=(2,2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dropout(0.5),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])

cnn.compile(
    loss='categorical_crossentropy',
    optimizer='adam',
    metrics=['accuracy']
)

cnn.summary()

history_cnn = cnn.fit(
    train_data,
    epochs = epochs,
    steps_per_epoch = 15,
    batch_size = batch_size,
    validation_data = test_data
)
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
conv2d_4 (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d_4 (MaxPooling 2D)	(None, 111, 111, 32)	0
conv2d_5 (Conv2D)	(None, 109, 109, 32)	9248
max_pooling2d_5 (MaxPooling 2D)	(None, 54, 54, 32)	0
flatten_4 (Flatten)	(None, 93312)	0
dropout_6 (Dropout)	(None, 93312)	0
dense_8 (Dense)	(None, 4)	373252
=====		
Total params: 383,396		
Trainable params: 383,396		
Non-trainable params: 0		

Epoch 1/10

```

15/15 [=====] - 86s 6s/step - loss: 299.4039 - accuracy: 0.3033 - val_loss: 6.7360 - val_accuracy: 0.3978
Epoch 2/10
15/15 [=====] - 87s 6s/step - loss: 4.1345 - accuracy: 0.4667 - val_loss: 1.8878 - val_accuracy: 0.4414
Epoch 3/10
15/15 [=====] - 85s 6s/step - loss: 1.8325 - accuracy: 0.4267 - val_loss: 1.4062 - val_accuracy: 0.4800
Epoch 4/10
15/15 [=====] - 77s 5s/step - loss: 1.4700 - accuracy: 0.4800 - val_loss: 1.2492 - val_accuracy: 0.5305
Epoch 5/10
15/15 [=====] - 87s 6s/step - loss: 1.2388 - accuracy: 0.5333 - val_loss: 1.1473 - val_accuracy: 0.5617
Epoch 6/10
15/15 [=====] - 86s 6s/step - loss: 1.1681 - accuracy: 0.5450 - val_loss: 1.0635 - val_accuracy: 0.5623
Epoch 7/10
15/15 [=====] - 76s 5s/step - loss: 1.1208 - accuracy: 0.5650 - val_loss: 1.0169 - val_accuracy: 0.5935
Epoch 8/10
15/15 [=====] - 89s 6s/step - loss: 1.0622 - accuracy: 0.6133 - val_loss: 0.9411 - val_accuracy: 0.6353
Epoch 9/10
15/15 [=====] - 84s 6s/step - loss: 1.1167 - accuracy: 0.5433 - val_loss: 0.9737 - val_accuracy: 0.6085
Epoch 10/10
15/15 [=====] - 89s 6s/step - loss: 0.9859 - accuracy: 0.5917 - val_loss: 0.9077 - val_accuracy: 0.6577


```

```

In [36]: plt.plot(history_cnn.history['val_accuracy'])
plt.plot(history_cnn.history['accuracy'])
plt.title('CNN Accuracy')
plt.legend(['Train', 'Test'])
plt.show()

```





## Pre-trained Model

```
In [34]: pre_trained = InceptionV3(weights = 'imagenet', classes = num_classes, inclu

for layer in pre_trained.layers:
    layer.trainable = False

inception = tf.keras.models.Sequential([
    tf.keras.layers.experimental.preprocessing.Rescaling(1./255, input_shape
pre_trained,
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(num_classes, activation='softmax')
])

inception.compile(
    loss = 'categorical_crossentropy',
    optimizer = 'adam',
    metrics = ['accuracy']
)

inception.summary()

history_inception = inception.fit(
    train_data,
    epochs = epochs,
    steps_per_epoch = 15,
    batch_size = batch_size,
    validation_data = test_data
)
```

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 224, 224, 3)	0
inception_v3 (Functional)	(None, 5, 5, 2048)	21802784
flatten_5 (Flatten)	(None, 51200)	0
dense_9 (Dense)	(None, 512)	26214912
batch_normalization_282 (Batch Normalization)	(None, 512)	2048
dropout_7 (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 4)	2052
Total params: 48,021,796		
Trainable params: 26,217,988		

Non-trainable params: 21,803,808

---

```


Epoch 1/10
15/15 [=====] - 303s 20s/step - loss: 1.6099 - accu
racy: 0.5967 - val_loss: 4.9819 - val_accuracy: 0.6072
Epoch 2/10
15/15 [=====] - 284s 20s/step - loss: 1.0382 - accu
racy: 0.6883 - val_loss: 1.8306 - val_accuracy: 0.6808
Epoch 3/10
15/15 [=====] - 286s 20s/step - loss: 0.7696 - accu
racy: 0.7050 - val_loss: 1.0974 - val_accuracy: 0.7263
Epoch 4/10
15/15 [=====] - 285s 20s/step - loss: 0.6924 - accu
racy: 0.7200 - val_loss: 1.1951 - val_accuracy: 0.6933
Epoch 5/10
15/15 [=====] - 285s 20s/step - loss: 0.7110 - accu
racy: 0.7200 - val_loss: 0.7085 - val_accuracy: 0.7537
Epoch 6/10
15/15 [=====] - 281s 20s/step - loss: 0.6277 - accu
racy: 0.7448 - val_loss: 0.7235 - val_accuracy: 0.7506
Epoch 7/10
15/15 [=====] - 285s 20s/step - loss: 0.5978 - accu
racy: 0.7700 - val_loss: 0.6816 - val_accuracy: 0.7618
Epoch 8/10
15/15 [=====] - 284s 20s/step - loss: 0.5678 - accu
racy: 0.7683 - val_loss: 0.5949 - val_accuracy: 0.7799
Epoch 9/10
15/15 [=====] - 284s 20s/step - loss: 0.6033 - accu
racy: 0.7633 - val_loss: 0.5736 - val_accuracy: 0.7787
Epoch 10/10
15/15 [=====] - 281s 20s/step - loss: 0.5254 - accu
racy: 0.7900 - val_loss: 0.7932 - val_accuracy: 0.7375

```

```

In [35]: plt.plot(history_inception.history['val_accuracy'])
plt.plot(history_inception.history['accuracy'])
plt.title('Inception Accuracy')
plt.legend(['Train', 'Test'])
plt.show()

```



I was going to take a picture of one of the playing cards I have at home, and see if the model would accurately classify it, but I ran out of time.

## Analysis

The 3 models used have very different levels of accuracy:

- Normal Sequential: ~25% accuracy
- CNN: ~59% accuracy
- Pre-Trained Model: ~79% accuracy

So just by glance, it would seem the Pre-Trained Model is the best, especially getting that accurate with the relatively small dataset given. However it also took significantly longer to train compared to the others. In general it seems that the cost of getting a more accurate model, tends to be an increase in training time.